



Unikernels: Paths to Production & Current Research Trends

Hugo Lefeuvre
The University of Manchester

ASPLOS 2022 Unikraft Tutorial, *March 1st*



A Decade of Unikernels...

~2012: A. Kantee's thesis on Rump kernels

A tad later, Mirage (ASPLOS'13), coined the term "unikernel".

A Decade of Unikernels...

~2012: A. Kantee's thesis on Rump kernels

A tad later, Mirage (ASPLOS'13), coined the term "unikernel".

Now, quite a large body of work

A Decade of Unikernels...

~2012: A. Kantee's thesis on Rump kernels

A tad later, Mirage (ASPLOS'13), coined the term "unikernel".

Now, quite a large body of work

ClickOS (NSDI'24) Mirage (ASPLOS'13)

Specialization / Performance

Unikraft (EuroSys'21)

A Decade of Unikernels...

~2012: A. Kantee's thesis on Rump kernels

A tad later, Mirage (ASPLOS'13), coined the term "unikernel".

Now, quite a large body of work

ClickOS (NSDI'24)

Mirage (ASPLOS'13)

Specialization / Performance

Unikraft (EuroSys'21)

LightVM (SOSP'17)

Firecracker (NSDI'21)

USETL (ApSys)

Scalability

Solo5 (HotCloud/SoCC)

A Decade of Unikernels...

~2012: A. Kantee's thesis on Rump kernels

A tad later, Mirage (ASPLOS'13), coined the term "unikernel".

Now, quite a large body of work

ClickOS (NSDI'24)

Mirage (ASPLOS'13)

Specialization / Performance

Unikraft (EuroSys'21)

LightVM (SOSP'17)

Firecracker (NSDI'21)

USETL (ApSys)

Scalability

Solo5 (HotCloud/SoCC)

EbbRT (OSDI'16)

Lupine Linux (EuroSys'20)

Compatibility

OSv (ATC'14)

HermiTux (VEE'19)

A Decade of Unikernels...

So, a lot of interest, a lot of nice ideas.

A Decade of Unikernels...

So, a lot of interest, a lot of nice ideas.

And *yet*, in practice, we don't see unikernels in production (even in the cloud) in 2022.

A Decade of Unikernels...

So, a lot of interest, a lot of nice ideas.

And *yet*, in practice, we don't see unikernels in production (even in the cloud) in 2022.

Why?

A Decade of Unikernels...

So, a lot of interest, a lot of nice ideas.

And *yet*, in practice, we don't see unikernels in production (even in the cloud) in 2022.

Why?

Current unikernels are just (mostly academic) **research prototypes**

Unikernels are Research Prototypes

Since they are research prototypes, none **managed (or even tried!) to:**

Unikernels are Research Prototypes

Since they are research prototypes, none **managed (or even tried!) to:**

- provide good compatibility
- excellent performance
- transparent integration in major deployment workflows

Unikernels are Research Prototypes

Since they are research prototypes, none **managed (or even tried!) to:**

- provide good compatibility
- excellent performance
- transparent integration in major deployment workflows
- take security seriously
- meet production-grade testing standards
- good debuggability

Unikernels are Research Prototypes

Since they are research prototypes, none **managed (or even tried!) to:**

- provide good compatibility
- excellent performance
- transparent integration in major deployment workflows
- take security seriously
- meet production-grade testing standards
- good debuggability

Now, if we want to see unikernels in production one day, one project must show that it can be done because some of these items are not trivial. **We want this project to be Unikraft.**

Unikernels are Research Prototypes

Since they are research prototypes, none **managed (or even tried!) to:**

- provide good compatibility
- excellent performance
- transparent integration in major deployment workflows
- take security seriously
- meet production-grade testing standards
- good debuggability

Now, if we want to see unikernels in production one day, one project must show that it can be done because some of these items are not trivial. **We want this project to be Unikraft.**

Unikraft is not there yet. But it's progressing, and we hope to see it reaching full maturity in the coming year

Unikernels are Research Prototypes

Since they are research prototypes, none **managed (or even tried!) to:**

- provide good compatibility →
- excellent performance →
- • transparent integration in major deployment workflows →
- take security seriously
- meet production-grade testing standards
- good debuggability

Now, if we want to see unikernels in production one day, one project must show that it can be done because some of these items are not trivial. **We want this project to be Unikraft.**

Unikraft is not there yet. But it's progressing, and we hope to see it reaching full maturity in the coming year

In the Works for Unikraft




Integration and Scalability



Security and Stability

In the Works for Unikraft




- Making Unikraft fit for classical deployment workflows (Kubernetes, etc.)

Integration and Scalability



Security and Stability

In the Works for Unikraft



Integration and Scalability



Security and Stability

- Making Unikraft fit for classical deployment workflows (Kubernetes, etc.)
- Making Unikraft just as debuggable as any userland application

In the Works for Unikraft




Integration and Scalability

- Making Unikraft fit for classical deployment workflows (Kubernetes, etc.)
- Making Unikraft just as debuggable as any userland application
- Open question: real world high-density: achieving 100s of unikernels per host?



Security and Stability

In the Works for Unikraft




Integration and Scalability



Security and Stability

- Making Unikraft fit for classical deployment workflows (Kubernetes, etc.)
- Making Unikraft just as debuggable as any userland application
- Open question: real world high-density: achieving 100s of unikernels per host?
- Matching the security/hardening features of mainstream OSes

In the Works for Unikraft



Integration and Scalability


- Making Unikraft fit for classical deployment workflows (Kubernetes, etc.)
- Making Unikraft just as debuggable as any userland application
- Open question: real world high-density: achieving 100s of unikernels per host?



Security and Stability

- Matching the security/hardening features of mainstream OSes
- Production-grade testing and fuzzing of Unikraft

In the Works for Unikraft




Integration and Scalability



Security and Stability

- Making Unikraft fit for classical deployment workflows (Kubernetes, etc.)
 - Making Unikraft just as debuggable as any userland application
 - Open question: real world high-density: achieving 100s of unikernels per host?
-
- Matching the security/hardening features of mainstream OSes
 - Production-grade testing and fuzzing of Unikraft
 - Specialization for the masses: automatic reasoning about Unikraft configurations

In the Works for Unikraft



Integration and Scalability



Security and Stability

- Making Unikraft fit for classical deployment workflows (Kubernetes, etc.)
 - Making Unikraft just as debuggable as any userland application
 - Open question: real world high-density: achieving 100s of unikernels per host?
-
- Matching the security/hardening features of mainstream OSes
 - Production-grade testing and fuzzing of Unikraft
 - Specialization for the masses: automatic reasoning about Unikraft configurations
 - Beyond the single trust domain: compartmentalizing Unikraft?

Integration & Scalability

Integration and Scalability

Why no unikernels in production?

[-] ▲ 🤖 yonkeltron 13 days ago | link

2

It's always been a mystery to me why Unikernels haven't caught on more. Especially with earlier toolkits like UniK and continuing work such as OSv. Does anyone have production experience or a tale of why they *didn't* pick Unikernels?

[-] ▲ 🌸 david_chisnall 13 days ago | link

5

The problems are not technical, they're economic. If I want to deploy a unikernel in the cloud, I am deploying as IaaS (i.e. a VM) where the unit of accounting is typically pairs of vCPUs and gigabytes of RAM, on an hourly basis. If I have the kind of problem where a unikernel would be a good solution, then I can deploy it as a FaaS system and be billed per CPU second and per RAM MiB second. **None of the cloud providers (yet?) have a way of deploying unikernels with FaaS-like pricing** and so if you make something small and efficient as a unikernel then it will have a load of unused CPU time and RAM that you're still being charged for. Unikernels only make economic sense if you're deploying your own datacenter.

No good integration.

https://lobste.rs/s/cyyx7a/unikraft_fast_secure_open_source (NOT an official Microsoft comment)

Integration and Scalability

Why no unikernels in production?

The screenshot shows a GitHub discussion thread titled "Why no unikernels in production?". The first comment by yonkeltron asks why Unikernels haven't caught on more, mentioning tools like UniK and OSv. The second comment by david_chisnall explains that the problems are economic, not technical, because cloud providers charge for VMs in pairs of vCPUs and RAM, making it difficult to deploy small unikernels. Below the discussion is a New Stack article titled "Unikernels Can't be Debugged, Joyent's Chief of Technology Argues", dated 25 Jan 2016, by Joab Jackson. The article is categorized under "CONTAINERS / SECURITY".

[-] ▲ yonkeltron 13 days ago | link
2
It's always been a mystery to me why Unikernels haven't caught on more. Especially with earlier toolkits like UniK and continuing work such as OSv. Does anyone have production experience or a tale of why they *didn't* pick Unikernels?

[-] ▲ david_chisnall 13 days ago | link
5
The problems are not technical, they're economic. If I want to deploy a unikernel in the cloud, I am deploying as IaaS (i.e. a VM) where the unit of accounting is typically pairs of vCPUs and gigabytes of RAM, on an hourly basis. If I would be a good solution, then I can deploy it as a FaaS (i.e. a function) where the unit of accounting is typically a few MiB second. None of the cloud providers (yet?) have a way to charge for a unikernel if you make something small and efficient as a unikernel that you're still being charged for. Unikernels only make sense if you can make them small enough to fit in a single vCPU.

THE NEW STACK Podcasts Events Ebooks Newsletter Sponsorship
Architecture Development Operations

CONTAINERS / SECURITY
Unikernels Can't be Debugged, Joyent's Chief of Technology Argues
25 Jan 2016 9:00am, by Joab Jackson

No good integration.

No good debugging facilities.

https://lobste.rs/s/cyyx7a/unikraft_fast_secure_open_source (NOT an official Microsoft comment)

<https://thenewstack.io/good-luck-debugging-unikernels-joyents-chief-technology-says/>

In the Works: Integration

People need integration of unikernels into orchestration frameworks to truly leverage their benefits.

In the Works: Integration

People need integration of unikernels into orchestration frameworks to truly leverage their benefits.

- Dynamically and quickly provision new services
- Schedule/Reschedule services based on workload (etc.)



In the Works: Integration

People need integration of unikernels into orchestration frameworks to truly leverage their benefits.

- Dynamically and quickly provision new services
- Schedule/Reschedule services based on workload (etc.)

No need to reinvent the wheel: make unikernels fit in these frameworks



In the Works: Integration

People need integration of unikernels into orchestration frameworks to truly leverage their benefits.

- Dynamically and quickly provision new services
- Schedule/Reschedule services based on workload (etc.)

No need to reinvent the wheel: make unikernels fit in these frameworks



We are almost there.

- Integration of unikernels in Kubernetes infrastucture
 - OCI-compliant unikernel runtime interface

In the Works: Integration

People need integration of unikernels into orchestration frameworks to truly leverage their benefits.

- Dynamically and quickly provision new services
- Schedule/Reschedule services based on workload (etc.)

No need to reinvent the wheel: make unikernels fit in these frameworks



We are almost there.

- Integration of unikernels in Kubernetes infrastucture
 - OCI-compliant unikernel runtime interface
- "A Unikernel in OCI Clothing": make unikernels look and feel like containers

A. Jung @ CNCF'21 https://www.youtube.com/watch?v=cV-xawN9_cg

In the Works: Integration

Good Progress

People need integration of unikernels into orchestration frameworks to truly leverage their benefits.

- Dynamically and quickly provision new services
- Schedule/Reschedule services based on workload (etc.)

No need to reinvent the wheel: make unikernels fit in these frameworks




We are almost there.

- Integration of unikernels in Kubernetes infrastucture
 - OCI-compliant unikernel runtime interface
- "A Unikernel in OCI Clothing": make unikernels look and feel like containers
- More progress to make on the FaaS side?

A. Jung @ CNCF'21 https://www.youtube.com/watch?v=cV-xawN9_cg


In the Works: Debugging



Vast engineering effort towards seamless introspection and debugging

S. Kuenzer, M. Rittinghaus @ FOSDEM'22 <https://fosdem.org/2022/schedule/event/skuenzer/>

In the Works: Debugging




Vast engineering effort towards seamless introspection and debugging

- Production-grade monitoring with Prometheus:
 - Monitor unikernels like any general-purpose VM
 - Setup alarms when values pass thresholds, etc.



S. Kuenzer, M. Rittinghaus @ FOSDEM'22 <https://fosdem.org/2022/schedule/event/skuenzer/>

In the Works: Debugging



Vast engineering effort towards seamless introspection and debugging


- Production-grade monitoring with Prometheus:
 - Monitor unikernels like any general-purpose VM
 - Setup alarms when values pass thresholds, etc.
- OS-level native GDB debugger support:
 - Debug unikernels like any userland application
 - Support for threads, OS-specific constructs (asserts, kernel crashes, etc.)
- Uniform debugging experience all platforms



S. Kuenzer, M. Rittinghaus @ FOSDEM'22 <https://fosdem.org/2022/schedule/event/skuenzer/>

In the Works: Debugging

Good Progress



Vast engineering effort towards seamless introspection and debugging

- Production-grade monitoring with Prometheus:
 - Monitor unikernels like any general-purpose VM
 - Setup alarms when values pass thresholds, etc.
- OS-level native GDB debugger support:
 - Debug unikernels like any userland application
 - Support for threads, OS-specific constructs (asserts, kernel crashes, etc.)
- Uniform debugging experience all platforms



S. Kuenzer, M. Rittinghaus @ FOSDEM'22 <https://fosdem.org/2022/schedule/event/skuenzer/>

Outstanding Question: Scalability?

Outstanding Question: Scalability?


As unikernels are evolving towards production, more challenges are likely to arise

We (as a community) claimed **massive packing** of such VMs on a single host

Outstanding Question: Scalability?

As unikernels are evolving towards production, more challenges are likely to arise

We (as a community) claimed **massive packing** of such VMs on a single host




LightVM (SOSP'17) paper: 1000s of noop unikernels on a single host

Outstanding Question: Scalability?

As unikernels are evolving towards production, more challenges are likely to arise

We (as a community) claimed **massive packing** of such VMs on a single host

... but **are we ready to see 50 Nginx instances on a single host?** Let alone 100s?



LightVM (SOSP'17) paper: 1000s of noop unikernels on a single host


Outstanding Question: Scalability?

As unikernels are evolving towards production, more challenges are likely to arise

We (as a community) claimed **massive packing** of such VMs on a single host

... but **are we ready to see 50 Nginx instances on a single host?** Let alone 100s?

Pretty much all unikernel papers evaluate systems with 1 CPU = 1vCPU static pinning...



LightVM (SOSP'17) paper: 1000s of noop unikernels on a single host

Outstanding Question: Scalability?


As unikernels are evolving towards production, more challenges are likely to arise

We (as a community) claimed **massive packing** of such VMs on a single host

... but **are we ready to see 50 Nginx instances on a single host?** Let alone 100s?

Pretty much all unikernel papers evaluate systems with 1 CPU = 1vCPU static pinning...

With such density: how do things look on the networking side? Have hypervisors really been thought for this kind of usage?



LightVM (SOSP'17) paper: 1000s of noop unikernels on a single host

Security & Stability

Security and Stability

Why no unikernels in production?

Because of security.

11 Conclusion



Much to the contrary of grandiose security claims often made by unikernel developers, the evidence thus far indicates that unikernels are decidedly *not* secure. [Bue] Having examined two major unikernels, Rumprun and IncludeOS, a worrying trend is already apparent: unikernels often lack even the most basic security features, especially with regard to memory corruption. ASLR, consistent W^X policy, and stack, heap, and standard library hardening are generally either missing, improperly implemented, or intentionally disabled. This would be bad enough in a full, general-purpose operating system, but it is made even worse in unikernels, where application and kernel code run together and share an address space. An attacker who gains code execution in the application can immediately go on to invoke kernel-level functionality, make hypercalls, perform raw packet I/O, and so on. This makes unikernels a particular liability when running alongside other types of hosts, as they can be used as pivot points from which to attack their neighbors with even more potency than would be possible on a full-OS VM or container (at least without privilege escalation).

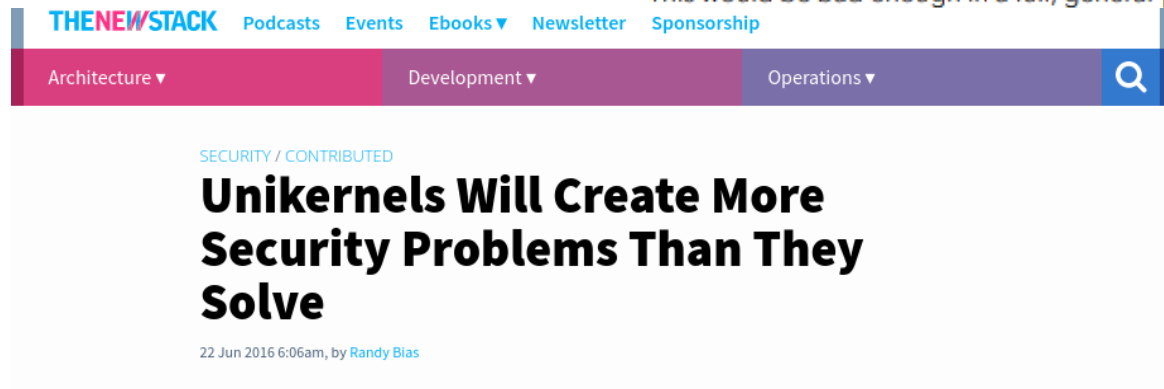
Given how low the bar has been set, there are numerous ways in which the currently abysmal state of unikernel security could improve. Aside from the protections we tested for – i.e. those typically found in modern, full-featured operating systems – there are several hypervisor-specific features that can be taken advantage of in order to improve unikernel security. For instance, many privileged operations, e.g. page table management, packet I/O, etc. can be performed via requests to the hypervisor rather than directly by the guest itself through emulated devices; such functionality is akin to syscalls or ioctls in a full OS.

Nonetheless, as it stands, unikernels remain an unsuitable and unappealing choice for production use, and will likely remain so until their security measures are at least brought in line with those of modern, full-featured operating systems.

https://www.nccgroup.com/globalassets/our-research/us/whitepapers/2019/ncc_group-assessing_unikernel_security.pdf

Security and Stability

Why no unikernels in production?



...and because of security 😊

11 Conclusion

Because of security.



Much to the contrary of grandiose security claims often made by unikernel developers, the evidence thus far indicates that unikernels are decidedly *not* secure. [Bue] Having examined two major unikernels, Rumprun and IncludeOS, a worrying trend is already apparent: unikernels often lack even the most basic security features, especially with regard to memory corruption. ASLR, consistent W^X policy, and stack, heap, and standard library hardening are generally either missing, improperly implemented, or intentionally disabled. This would be bad enough in a full, general-purpose operating system, but it is made even worse in uniker-

together and share an address space. An attacker who gains access to the guest immediately goes on to invoke kernel-level functionality, make hypercalls, and makes unikernels a particular liability when running alongside other VMs or containers (at least without privilege escalation).

There are numerous ways in which the currently abysmal state of unikernel security is the result of the protections we tested for – i.e. those typically found in full operating systems. There are several hypervisor-specific features that can be taken advantage of to improve security. For instance, many privileged operations, e.g. page faults, are performed via requests to the hypervisor rather than directly by the guest itself through emulated devices; such functionality is akin to syscalls or ioctls in a full OS.

Nonetheless, as it stands, unikernels remain an unsuitable and unappealing choice for production use, and will likely remain so until their security measures are at least brought in line with those of modern, full-featured operating systems.

https://www.nccgroup.com/globalassets/our-research/us/whitepapers/2019/ncc_group-assessing_unikernel_security.pdf
<https://thenewstack.io/unikernels-will-create-security-problems-solve/>

Getting in Line with Mainstream OSes

Getting in Line with Mainstream OSes

Write or Execute

The screenshot shows a GitHub pull request page for the repository 'unikraft / unikraft'. The pull request is titled 'Virtual Memory API (x86_64, kvm) #338' and is in the 'Open' state. It was created by 'marcrittinghaus' and wants to merge 11 commits into the 'unikraft:staging' branch from the 'marcrittinghaus:mritting/paging' branch. The page shows 7 conversations, 11 commits, 0 checks, and 41 files changed. A comment from 'marcrittinghaus' dated Nov 22, 2021, contains a 'Prerequisite checklist' with four items: reading contribution guidelines, testing changes, running 'checkpatch.pl', and updating documentation. Below the checklist is a 'Base target' section listing architecture (x86_64), platform (kvm), and application (N/A). An 'Additional configuration' section describes a new configuration option 'CONFIG_PAGING=y' and mentions dependencies on 'ukfalloc' and 'ukfallocbuddy'.

unikraft / unikraft Public

<> Code 93 Pull requests 51 Discussions Actions Projects 1 Security

Virtual Memory API (x86_64, kvm) #338

Open marcrittinghaus wants to merge 11 commits into unikraft:staging from marcrittinghaus:mritting/paging

Conversation 7 Commits 11 Checks 0 Files changed 41

marcrittinghaus commented on Nov 22, 2021

Prerequisite checklist

- ☒ Read the [contribution guidelines](#) regarding submitting new changes to the project;
- ☒ Tested your changes against relevant architectures and platforms;
- ☒ Ran the [checkpatch.pl](#) on your commit series before opening this PR;
- ☒ Updated relevant documentation.

Base target

- Architecture(s): x86_64
- Platform(s): kvm
- Application(s): N/A

Additional configuration

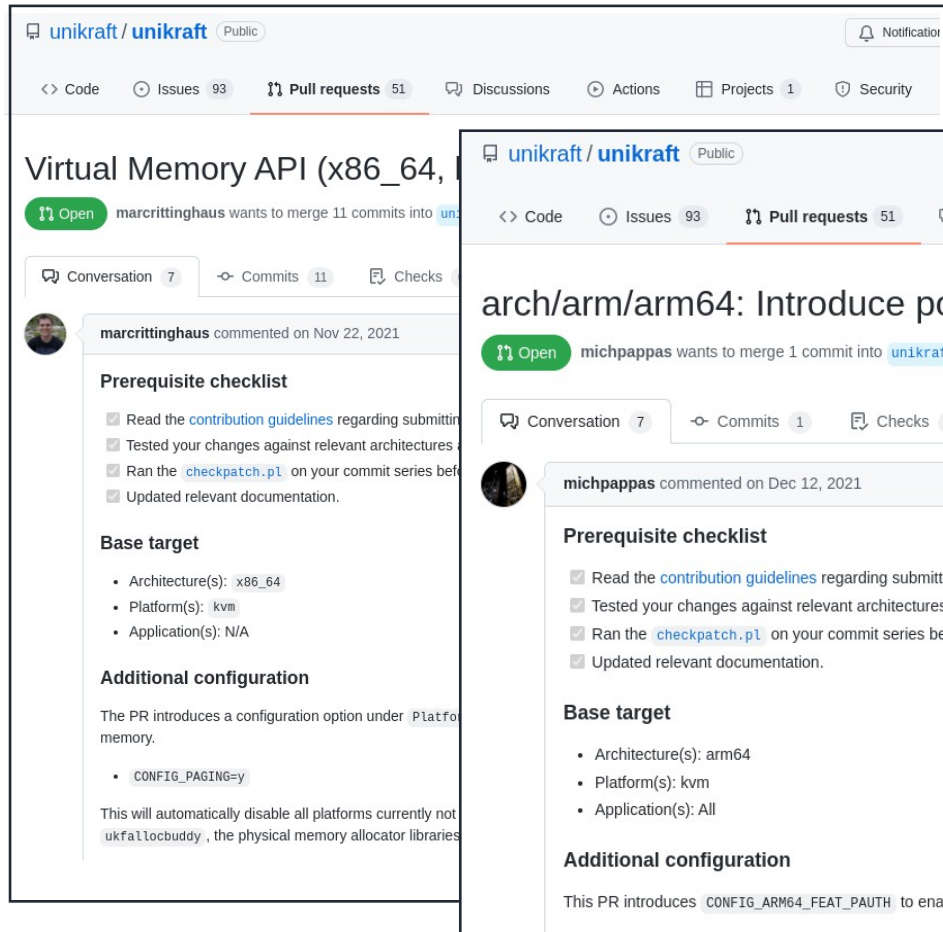
The PR introduces a configuration option under Platform Configuration | Platform Interface Options to enable virtual memory.

- CONFIG_PAGING=y

This will automatically disable all platforms currently not supported (e.g., Xen) and also add a dependency on ukfalloc and ukfallocbuddy, the physical memory allocator libraries.

Getting in Line with Mainstream OSes

Write or Execute



unikraft / unikraft Public

<> Code Issues 93 Pull requests 51 Discussions Actions Projects 1 Security

Virtual Memory API (x86_64, ...)

marcrittinghaus wants to merge 11 commits into unikraft:staging

Conversation 7 Commits 11 Checks

marcrittinghaus commented on Nov 22, 2021

Prerequisite checklist

- ☒ Read the [contribution guidelines](#) regarding submitting new changes to the project;
- ☒ Tested your changes against relevant architectures and platforms;
- ☒ Ran the [checkpatch.pl](#) on your commit series before opening this PR;
- ☒ Updated relevant documentation.

Base target

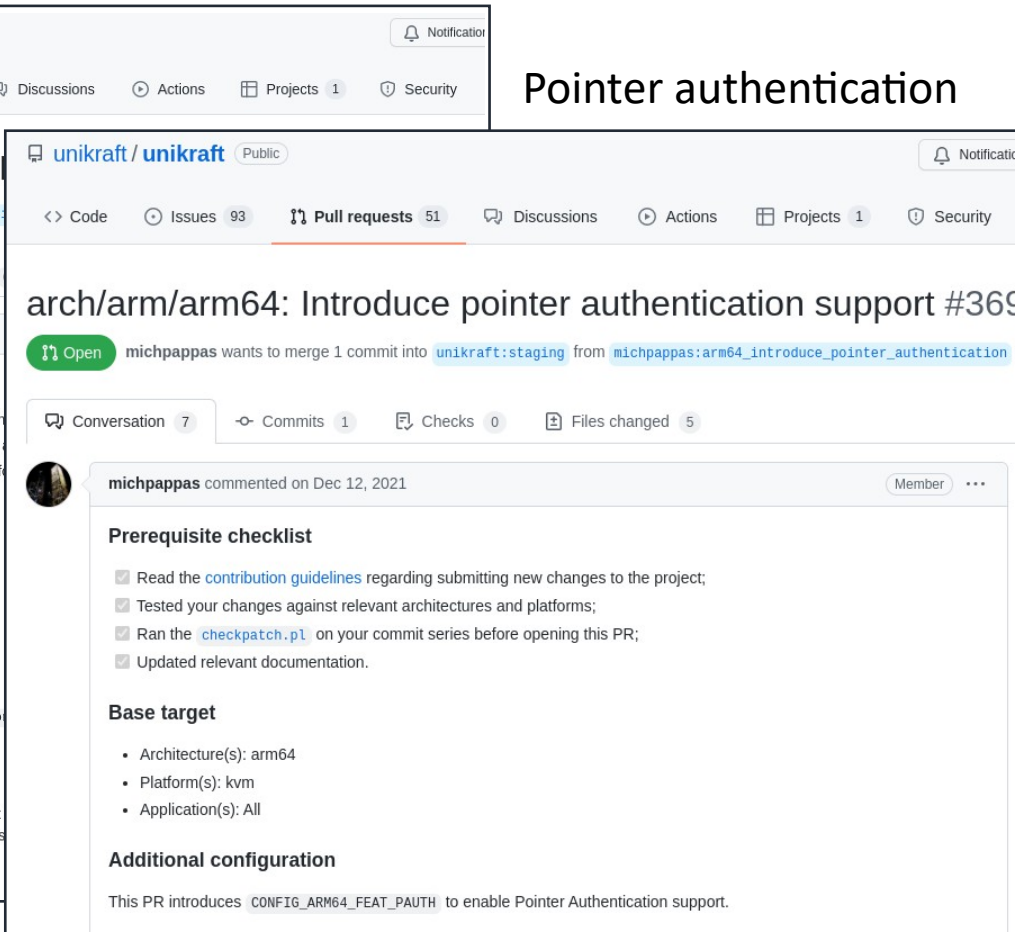
- Architecture(s): x86_64
- Platform(s): kvm
- Application(s): N/A

Additional configuration

The PR introduces a configuration option under Platform memory.

- CONFIG_PAGING=y

This will automatically disable all platforms currently not supported by ukfallocbuddy, the physical memory allocator libraries.



unikraft / unikraft Public

<> Code Issues 93 Pull requests 51 Discussions Actions Projects 1 Security

arch/arm/arm64: Introduce pointer authentication support #369

michpappas wants to merge 1 commit into unikraft:staging from michpappas:arm64_introduce_pointer_authentication

Conversation 7 Commits 1 Checks 0 Files changed 5

michpappas commented on Dec 12, 2021

Prerequisite checklist

- ☒ Read the [contribution guidelines](#) regarding submitting new changes to the project;
- ☒ Tested your changes against relevant architectures and platforms;
- ☒ Ran the [checkpatch.pl](#) on your commit series before opening this PR;
- ☒ Updated relevant documentation.

Base target

- Architecture(s): arm64
- Platform(s): kvm
- Application(s): All

Additional configuration

This PR introduces CONFIG_ARM64_FEAT_PAUTH to enable Pointer Authentication support.

Pointer authentication

Getting in Line with Mainstream OSes

Write or Execute

The image displays three overlapping GitHub pull request (PR) screenshots from the `unikraft/unikraft` repository, illustrating contributions to mainstream OSes.

Left Screenshot: Virtual Memory API (x86_64, ...)

- PR Title:** Virtual Memory API (x86_64, ...)
- Author:** marcrittinghaus
- Target:** `unikraft:staging`
- Prerequisite checklist:**
 - ☒ Read the [contribution guidelines](#) regarding submitting new changes to the project;
 - ☒ Tested your changes against relevant architectures and platforms;
 - ☒ Ran the `checkpatch.pl` on your commit series before opening this PR;
 - ☒ Updated relevant documentation.
- Base target:**
 - Architecture(s): `x86_64`
 - Platform(s): `kvm`
 - Application(s): `N/A`
- Additional configuration:**
 - `CONFIG_PAGING=y`

Middle Screenshot: arch/arm/arm64: Introduce pointer authentication

- PR Title:** arch/arm/arm64: Introduce pointer authentication
- Author:** michpappas
- Target:** `unikraft:staging`
- Prerequisite checklist:**
 - ☒ Read the [contribution guidelines](#) regarding submitting new changes to the project;
 - ☒ Tested your changes against relevant architectures and platforms;
 - ☒ Ran the `checkpatch.pl` on your commit series before opening this PR;
 - ☒ Updated relevant documentation.
- Base target:**
 - Architecture(s): `arm64`
 - Platform(s): `kvm`
 - Application(s): `All`
- Additional configuration:**
 - `CONFIG_ARM64_FEAT_PAUTH`

Right Screenshot: build: Option to compile as PIE #239

- PR Title:** build: Option to compile as PIE #239
- Author:** daniel20
- Target:** `unikraft:staging`
- Comment:**

This patch adds the option to compile the unikernel as a position-independent executable so we can have ASLR. If the unikernel is compiled as PIE, it cannot run on its own. A bootloader will be needed that will come in a future PR.
- Signed-off-by:** Daniel Dinca dincadaniel97@gmail.com

Getting in Line with Mainstream OSes

Write or Execute

... and stack protection, KASan, etc.

The image displays three overlapping GitHub pull request (PR) pages from the `unikraft/unikraft` repository, illustrating various security and system-level changes.

Leftmost PR: Virtual Memory API (x86_64, ...)

- Prerequisite checklist:**
 - ☒ Read the [contribution guidelines](#) regarding submitting new changes to the project;
 - ☒ Tested your changes against relevant architectures and platforms;
 - ☒ Ran the `checkpatch.pl` on your commit series before opening this PR;
 - ☒ Updated relevant documentation.
- Base target:**
 - Architecture(s): `x86_64`
 - Platform(s): `kvm`
 - Application(s): `N/A`
- Additional configuration:**

The PR introduces a configuration option under `Platform` memory.

 - `CONFIG_PAGING=y`

This will automatically disable all platforms currently not supporting `ukfallocbuddy`, the physical memory allocator libraries.

Middle PR: arch/arm/arm64: Introduce pointer authentication

- Prerequisite checklist:**
 - ☒ Read the [contribution guidelines](#) regarding submitting new changes to the project;
 - ☒ Tested your changes against relevant architectures and platforms;
 - ☒ Ran the `checkpatch.pl` on your commit series before opening this PR;
 - ☒ Updated relevant documentation.
- Base target:**
 - Architecture(s): `arm64`
 - Platform(s): `kvm`
 - Application(s): `All`
- Additional configuration:**

This PR introduces `CONFIG_ARM64_FEAT_PAUTH` to enable Pointer Authentication support.

Rightmost PR: build: Option to compile as PIE #239

- Comment by daniel20 (Jun 25, 2021):**

This patch adds the option to compile the unikernel as a position-independent executable so we can have ASLR. If the unikernel is compiled as PIE, it cannot run on its own. A bootloader will be needed that will come in a future PR.

Signed-off-by: Daniel Dinca dincadaniel97@gmail.com

Getting in Line with Mainstream OSes

Good Progress

Write or Execute

... and stack protection, KASan, etc.

unikraft / unikraft Public

<> Code Issues 93 Pull requests 51 Discussions Actions Projects 1 Security

Virtual Memory API (x86_64, ...)

marcrittinghaus wants to merge 11 commits into unikraft:staging from marcrittinghaus:vm-api

Conversation 7 Commits 11 Checks

marcrittinghaus commented on Nov 22, 2021

Prerequisite checklist

- ☒ Read the [contribution guidelines](#) regarding submitting new changes to the project;
- ☒ Tested your changes against relevant architectures and platforms;
- ☒ Ran the [checkpatch.pl](#) on your commit series before opening this PR;
- ☒ Updated relevant documentation.

Base target

- Architecture(s): x86_64
- Platform(s): kvm
- Application(s): N/A

Additional configuration

The PR introduces a configuration option under Platform memory.

- CONFIG_PAGING=y

This will automatically disable all platforms currently not supporting `ukfallocbuddy`, the physical memory allocator libraries.

Pointer authentication

unikraft / unikraft Public

<> Code Issues 93 Pull requests 51 Discussions Actions Projects 1 Security

arch/arm/arm64: Introduce pointer authentication

michpappas wants to merge 1 commit into unikraft:staging from michpappas:arm64_1

Conversation 7 Commits 1 Checks 0 Files changed 5

michpappas commented on Dec 12, 2021

Prerequisite checklist

- ☒ Read the [contribution guidelines](#) regarding submitting new changes to the project;
- ☒ Tested your changes against relevant architectures and platforms;
- ☒ Ran the [checkpatch.pl](#) on your commit series before opening this PR;
- ☒ Updated relevant documentation.

Base target

- Architecture(s): arm64
- Platform(s): kvm
- Application(s): All

Additional configuration

This PR introduces `CONFIG_ARM64_FEAT_PAUTH` to enable Pointer Authentication support.

ASLR

unikraft / unikraft Public

<> Code Issues 93 Pull requests 51 Discussions Actions Projects 1 Security

build: Option to compile as PIE #239

daniel20 wants to merge 1 commit into unikraft:staging from daniel20:ddinca/build-as-PIE

Conversation 7 Commits 1 Checks 0 Files changed 3

daniel20 commented on Jun 25, 2021

This patch adds the option to compile the unikernel as a position-independent executable so we can have ASLR. If the unikernel is compiled as PIE, it cannot run on it's own. A bootloader will be needed that will come in a future PR.

Signed-off-by: Daniel Dinca dincadaniel97@gmail.com

In the Works: Testing, Fuzz Testing

To go mainstream, unikernels need not only hardening that's in line with mainstream OSes, but also **production-grade testing**

In the Works: Testing, Fuzz Testing

To go mainstream, unikernels need not only hardening that's in line with mainstream OSes, but also **production-grade testing**

Production-grade testing =

- Continuous testing (CI/CD, test suite)
- Destructive testing (Fuzzing)

In the Works: Testing, Fuzz Testing

To go mainstream, unikernels need not only hardening that's in line with mainstream OSes, but also **production-grade testing**

Production-grade testing =

- Continuous testing (CI/CD, test suite)
- Destructive testing (Fuzzing)

Significant efforts on continuous testing:

In the Works: Testing, Fuzz Testing

Good Progress

To go mainstream, unikernels need not only hardening that's in line with mainstream OSes, but also **production-grade testing**

Production-grade testing =

- Continuous testing (CI/CD, test suite)
- Destructive testing (Fuzzing)



Significant efforts on continuous testing:

- CI/CD pipeline tests patches systematically (Concourse)
- Application-level tests but also kernel unit-tests (uktest)

A. Jung @ FOSDEM'22

https://fosdem.org/2022/schedule/event/massive_unikernel_matrices_with_unikraft_concourse_and_more/

In the Works: Testing, Fuzz Testing

To go mainstream, unikernels need not only hardening that's in line with mainstream OSes, but also **production-grade testing**

Ongoing efforts on fuzzing as well

In the Works: Testing, Fuzz Testing

To go mainstream, unikernels need not only hardening that's in line with mainstream OSes, but also **production-grade testing**

Ongoing efforts on fuzzing as well

- Not entirely trivial, as most OS fuzzers are **tailored for Linux**

In the Works: Testing, Fuzz Testing

To go mainstream, unikernels need not only hardening that's in line with mainstream OSes, but also **production-grade testing**

Ongoing efforts on fuzzing as well

- Not entirely trivial, as most OS fuzzers are **tailored for Linux**

Not "just" a matter
of porting Syzkaller
to Unikraft

In the Works: Testing, Fuzz Testing

To go mainstream, unikernels need not only hardening that's in line with mainstream OSes, but also **production-grade testing**

Ongoing efforts on fuzzing as well

- Not entirely trivial, as most OS fuzzers are **tailored for Linux**
 - Coverage measurement (no Kcov, porting to gcov not without changes)
 - Not every system call is fully implemented

Not "just" a matter of porting Syzkaller to Unikraft

In the Works: Testing, Fuzz Testing

To go mainstream, unikernels need not only hardening that's in line with mainstream OSes, but also **production-grade testing**

Ongoing efforts on fuzzing as well

- Not entirely trivial, as most OS fuzzers are **tailored for Linux**
 - Coverage measurement (no Kcov, porting to gcov not without changes)
 - Not every system call is fully implemented
- How does unikernel fuzzing impact the architecture of fuzzers?
- How to design a fuzzer that's **ready to "plug and play" in any POSIX OS?**

Not "just" a matter of porting Syzkaller to Unikraft

In the Works: Automatic Specialization

Another end of the "testing" topic: how do you determine **how good a configuration really is?**

A. Jung et al. @ APSys'21 <https://www.youtube.com/watch?v=YLf86gcHW4E>

In the Works: Automatic Specialization

Another end of the "testing" topic: how do you determine **how good a configuration really is?**

Unikernels pitch specialization:

- Your best Nginx configuration is not your best SQLite configuration
- ...and probably not your best Redis configuration either 😊

A. Jung et al. @ APSys'21 <https://www.youtube.com/watch?v=YLf86gcHW4E>

In the Works: Automatic Specialization

Another end of the "testing" topic: how do you determine **how good a configuration really is?**

Unikernels pitch specialization:

- Your best Nginx configuration is not your best SQLite configuration
- ...and probably not your best Redis configuration either 😊

The number of possible configurations: astronomical scale

A. Jung et al. @ APSys'21 <https://www.youtube.com/watch?v=YLf86gcHW4E>

In the Works: Automatic Specialization

Another end of the "testing" topic: how do you determine **how good a configuration really is?**

Unikernels pitch specialization:

- Your best Nginx configuration is not your best SQLite configuration
- ...and probably not your best Redis configuration either 😊

The number of possible configurations: astronomical scale

- Small subset of configuration options (Nginx) $\sim 10^{13}$
- How do you explore this? Can you use optimization algorithms? ML?

A. Jung et al. @ APSys'21 <https://www.youtube.com/watch?v=YLf86gcHW4E>

In the Works: Compartmentalization

In the Works: Compartmentalization

Traditional understanding of unikernels:

Unikernel = one single trust domain (kernel + application)

(Used to) make sense.

In the Works: Compartmentalization

Traditional understanding of unikernels:

Unikernel = one single trust domain (kernel + application)

(Used to) make sense.

Certain applications are large, with heterogeneous components: trust, safety, properties, requirements...

In the Works: Compartmentalization

Traditional understanding of unikernels:

Unikernel = one single trust domain (kernel + application)

(Used to) make sense.

Certain applications are large, with heterogeneous components: trust, safety, properties, requirements...

And at the same time we see (re-)appearing a lot of lightweight isolation mechanisms (protection keys, HW capabilities, SFI, etc.)

In the Works: Compartmentalization

There is an opportunity to use these mechanisms to make unikernels even safer without yielding their benefits!

In the Works: Compartmentalization

There is an opportunity to use these mechanisms to make unikernels even safer without yielding their benefits!



This is what initially motivated our work FlexOS: can we reconcile unikernels/libOSes with isolation to obtain a new OS model that offers not only specialization **towards performance, but also towards safety?**

H. Lefeuvre et al. @ ASPLOS'22, come to our talk Thursday morning!
(also @ FOSDEM'22 https://fosdem.org/2022/schedule/event/tee_flexos/)

In the Works: Compartmentalization

There is an opportunity to use these mechanisms to make unikernels even safer without yielding their benefits!



This is what initially motivated our work FlexOS: can we reconcile unikernels/libOSes with isolation to obtain a new OS model that offers not only specialization **towards performance, but also towards safety?**

Other groups explored this direction: CubicleOS (also ASPLOS, 2021). Explore intra-unikernel isolation with Intel MPK.

H. Lefeuve et al. @ ASPLOS'22, come to our talk Thursday morning!
(also @ FOSDEM'22 https://fosdem.org/2022/schedule/event/tee_flexos/)

How about Compatibility?

Compatibility: a Solved Problem?

Historical unikernels: hand-written for an application (ClickOS).

Compatibility: a Solved Problem?

Historical unikernels: hand-written for an application (ClickOS)

Terrible Compatibility

Compatibility: a Solved Problem?

Historical unikernels: hand-written for an application (ClickOS),

Terrible Compatibility

But we're not in the 2010s anymore, Unikraft aims at Linux/POSIX compliance.

Compatibility: a Solved Problem?

Historical unikernels: hand-written for an application (ClickOS),

Terrible Compatibility

But we're not in the 2010s anymore, Unikraft aims at Linux/POSIX compliance.

Growing number of supported system calls: now 170+

As a point of comparison, Graphene (Gramine) also supports about ~170

Compatibility: a Solved Problem?

Historical unikernels: hand-written for an application (ClickOS),

Terrible Compatibility

But we're not in the 2010s anymore, Unikraft aims at Linux/POSIX compliance.

Growing number of supported system calls: now 170+

As a point of comparison, Graphene (Gramine) also supports about ~170

To be clear: Unikraft is neither fully Linux compatible, nor fully POSIX compliant!

- The good old `fork()` problem
- Not all system calls are fully implemented

But **do you really need to be fully compatible** to be useful?

Compatibility: a Solved Problem?

But **do you really need to be fully compatible** to be useful?

Many applications that use unsupported features tend to be bad candidates for unikernelization anyways (system admin tools, heavily multiprocess apps that cannot use threads).

Compatibility: a Solved Problem?

But **do you really need to be fully compatible** to be useful?

Many applications that use unsupported features tend to be bad candidates for unikernelization anyways (system admin tools, heavily multiprocess apps that cannot use threads).

For the rest, partial compatibility is just fine if porting is a reasonable task.

And in 2022, it is.

Compatibility: a Solved Problem?

But **do you really need to be fully compatible** to be useful?

Many applications that use unsupported features tend to be bad candidates for unikernelization anyways (system admin tools, heavily multiprocess apps that cannot use threads).

For the rest, partial compatibility is just fine if porting is a reasonable task.

And in 2022, it is.

Given the benefits of Unikraft, a week of porting is a minor annoyance.
All you need is a good application test-suite (but you have one, right? 😊)

H. Lefeuve et al. @ USENIX ;login

<https://www.usenix.org/publications/loginonline/unikraft-and-coming-age-unikernels>

Conclusion

In a Nutshell

In the last decade we've seen a lot of unikernels come and go, but **none that managed to:**

In a Nutshell

In the last decade we've seen a lot of unikernels come and go, but **none that managed to:**

- provide good compatibility
- excellent performance
- transparent integration in major deployment workflows

In a Nutshell

In the last decade we've seen a lot of unikernels come and go, but **none that managed to:**

- provide good compatibility
- excellent performance
- transparent integration in major deployment workflows
- take security seriously
- meet production-grade testing standards
- good debuggability

In a Nutshell

In the last decade we've seen a lot of unikernels come and go, but **none that managed to:**

- provide good compatibility
- excellent performance
- transparent integration in major deployment workflows
- take security seriously
- meet production-grade testing standards
- good debuggability

And that's fine, because they were just **research projects**

In a Nutshell

In the last decade we've seen a lot of unikernels come and go, but **none that managed to:**

- provide good compatibility
- excellent performance
- transparent integration in major deployment workflows
- take security seriously
- meet production-grade testing standards
- good debuggability

And that's fine, because they were just **research projects**

Now, if we want to see unikernels in production one day, one project must manage it. **We want this project to be Unikraft.**

In a Nutshell

In the last decade we've seen a lot of unikernels come and go, but **none that managed to:**

- provide good compatibility
- excellent performance
- transparent integration in major deployment workflows
- take security seriously
- meet production-grade testing standards
- good debuggability

And that's fine, because they were just **research projects**

Now, if we want to see unikernels in production one day, one project must manage it. **We want this project to be Unikraft.**

Unikraft is not there yet. But it's progressing, and we hope to see it reaching full maturity in the coming year

In a Nutshell

Over time, doing all this engineering proved **fruitful on the research side**

In a Nutshell

Over time, doing all this engineering proved **fruitful on the research side**

Unikraft turns out to be an **excellent substrate for top-tier publications**

It's open, small, clean, modular, fast to experiment with, such that undergrads and grad students quickly get to understand it

In a Nutshell

Over time, doing all this engineering proved **fruitful on the research side**

Unikraft turns out to be an **excellent substrate for top-tier publications**

It's open, small, clean, modular, fast to experiment with, such that undergrads and grad students quickly get to understand it

Probably one of the good examples where starting clean-slate pays out in the long run

In a Nutshell

Over time, doing all this engineering proved **fruitful on the research side**

Unikraft turns out to be an **excellent substrate for top-tier publications**

It's open, small, clean, modular, fast to experiment with, such that undergrads and grad students quickly get to understand it

Probably one of the good examples where starting clean-slate pays out in the long run

What will the broader systems community build with Unikraft?



Pushing Unikernels to Production!

Unikraft Community: <https://unikraft.org/>

Unikraft Cloud: <https://unikraft.io/>

Code: <https://github.com/unikraft>